



Introduction to KERAS Framework

Why we need a framework?

Simple Neural Network



Deep Learning Neural Network



● Input Layer


● Hidden Layer

● Output Layer


Forward Computing

Why we need a framework?

Simple Neural Network



Deep Learning Neural Network



● Input Layer

● Hidden Layer


● Output Layer

Backward updating


Why We Need a Framework?

It involve a lot of Computation

Simple Neural Network



Deep Learning Neural Network



● Input Layer

● Hidden Layer

● Output Layer

Backward updating

Why We Need a Framework?


Even for a simple neural network:

- hundreds lines of code is needed, in python.
- In a DL framework (TensorFlow, PyTorch, etc) you may write it in 20-50 lines.
- In Keras, it can be done in 5 lines of code!

```
1 model = Sequential()  
2 model.add(Dense(5, activation='relu', input_shape=(784,)))  
3 model.add(Dense(10, activation='softmax'))  
4 model.compile(loss='categorical_crossentropy', optimizer='sgd')  
5 history = model.fit(x_train, y_train, epochs = 10)
```

What a DL framework usually do for us?

- Computes Automatic Backpropagation.
- Contains deep learning networks.
- makes computing on GPU easier.



Why we have chosen Keras?

Keras: Deep Learning for humans



Keras

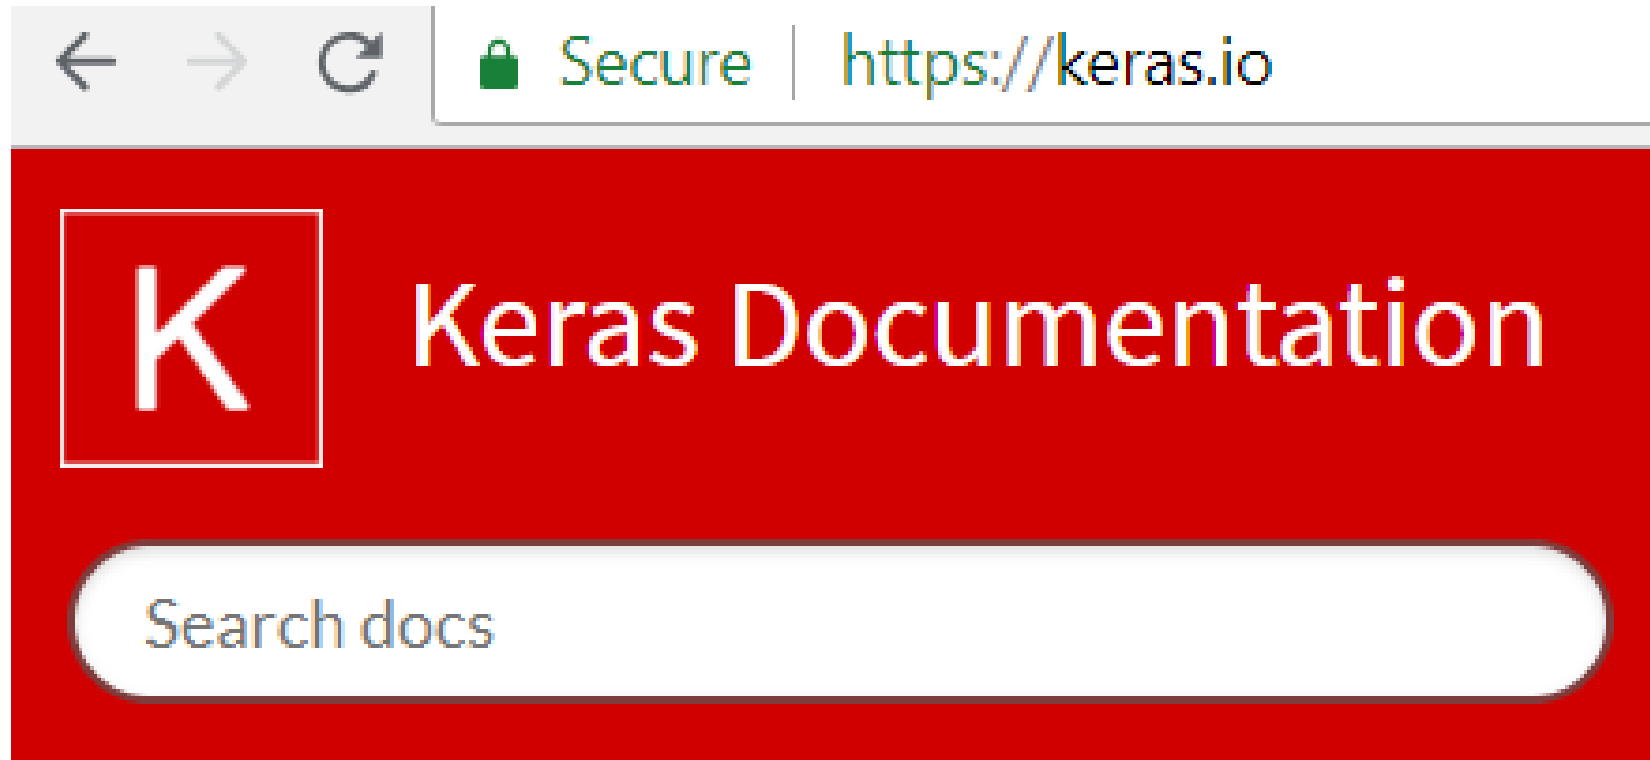
build passing license MIT

You have just found Keras.

- Allows for easy and fast prototyping
- Supports both convolutional and recurrent networks, as well as combinations of the two.
- Runs seamlessly on CPU and GPU.

What more about Keras?

- Good Documentation at: Keras.io



What more about Keras?

- You can use TensorFlow (generally any backend) code in Keras:

```
from keras import backend as K
```

What more about Keras?

- You can use Keras in TensorFlow

```
import tensorflow as tf  
tf.keras.
```

A deep learning problem in Keras?

- You can use Keras in TensorFlow



- keras.datasets
 - MNIST database of handwritten digits
 - CIFAR-10 small image classification
 - CIFAR-100 small image classification
 - IMDB Movie Reviews Sentiment Classification
 - ...

```
from keras.datasets import fashion_mnist

(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
```

- `keras.preprocessing.image`:
 - `ImageDataGenerator`:
 - Generate Batches of tensor image data
 - Feature-wise and sample-wise normalization
 - Whitening
 - Any other preprocessing function.
 - Train Validation Split.
 - Real-Time data Augmentation (will be covered in CNN section)

- `keras.preprocessing.text`:
 - `text_to_word_sequence`
 - `one_hot`
 - ...
- `keras.preprocessing.sequence`:
 - TimeseriesGenerator Class
 - `pad_sequences`

Data
Preprocessing


Defining
Model

Training

Evaluating


Utilities

Network
state



(Should be initialized)

Input



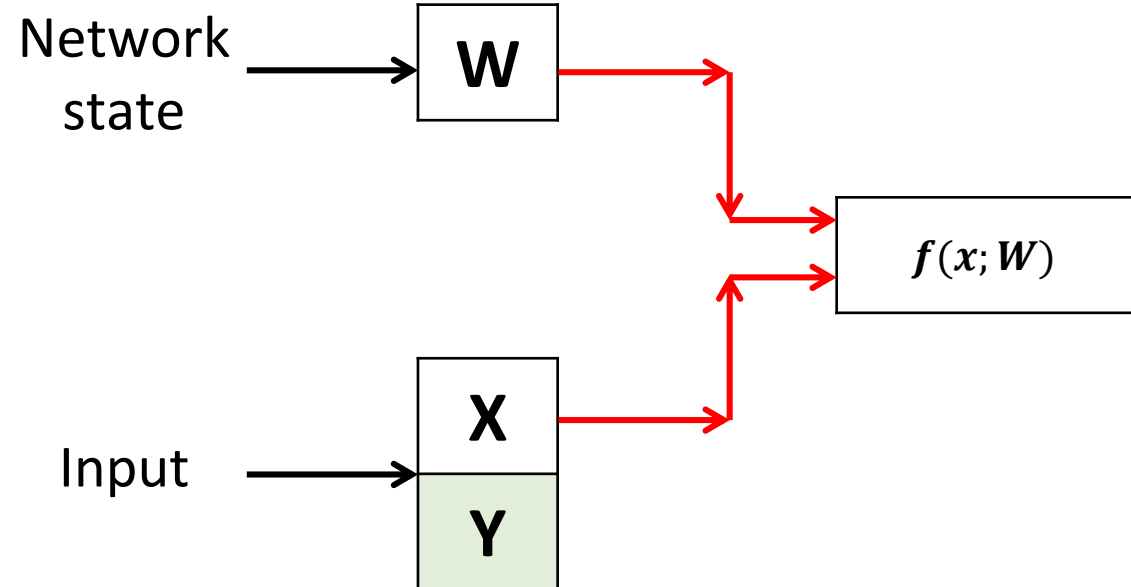
Data
Preprocessing


Defining
Model

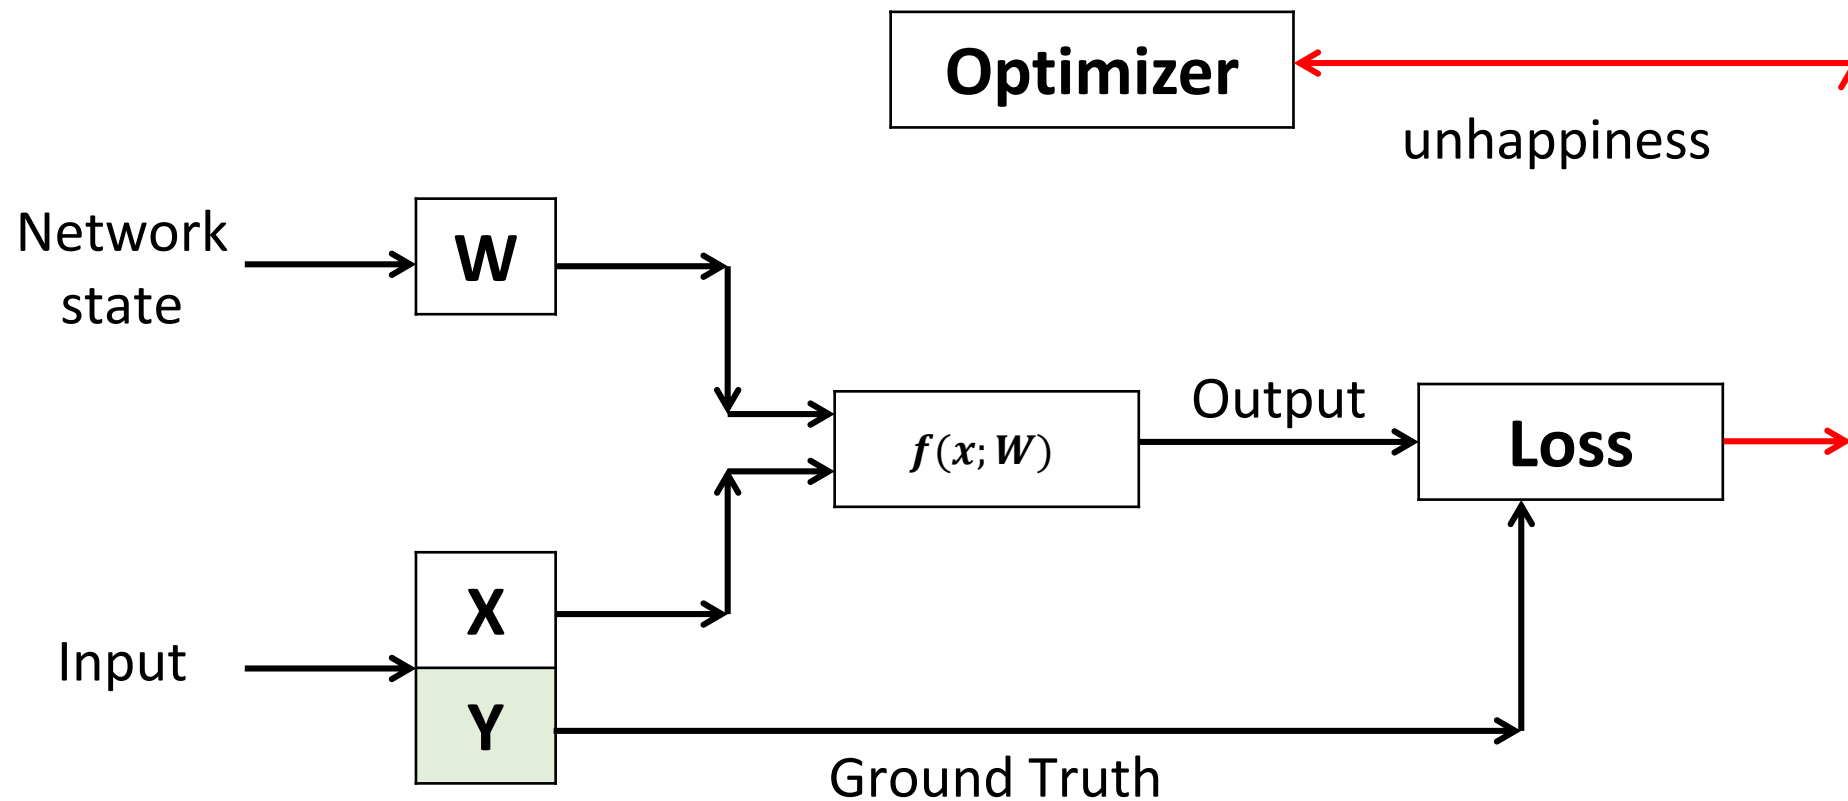
Training

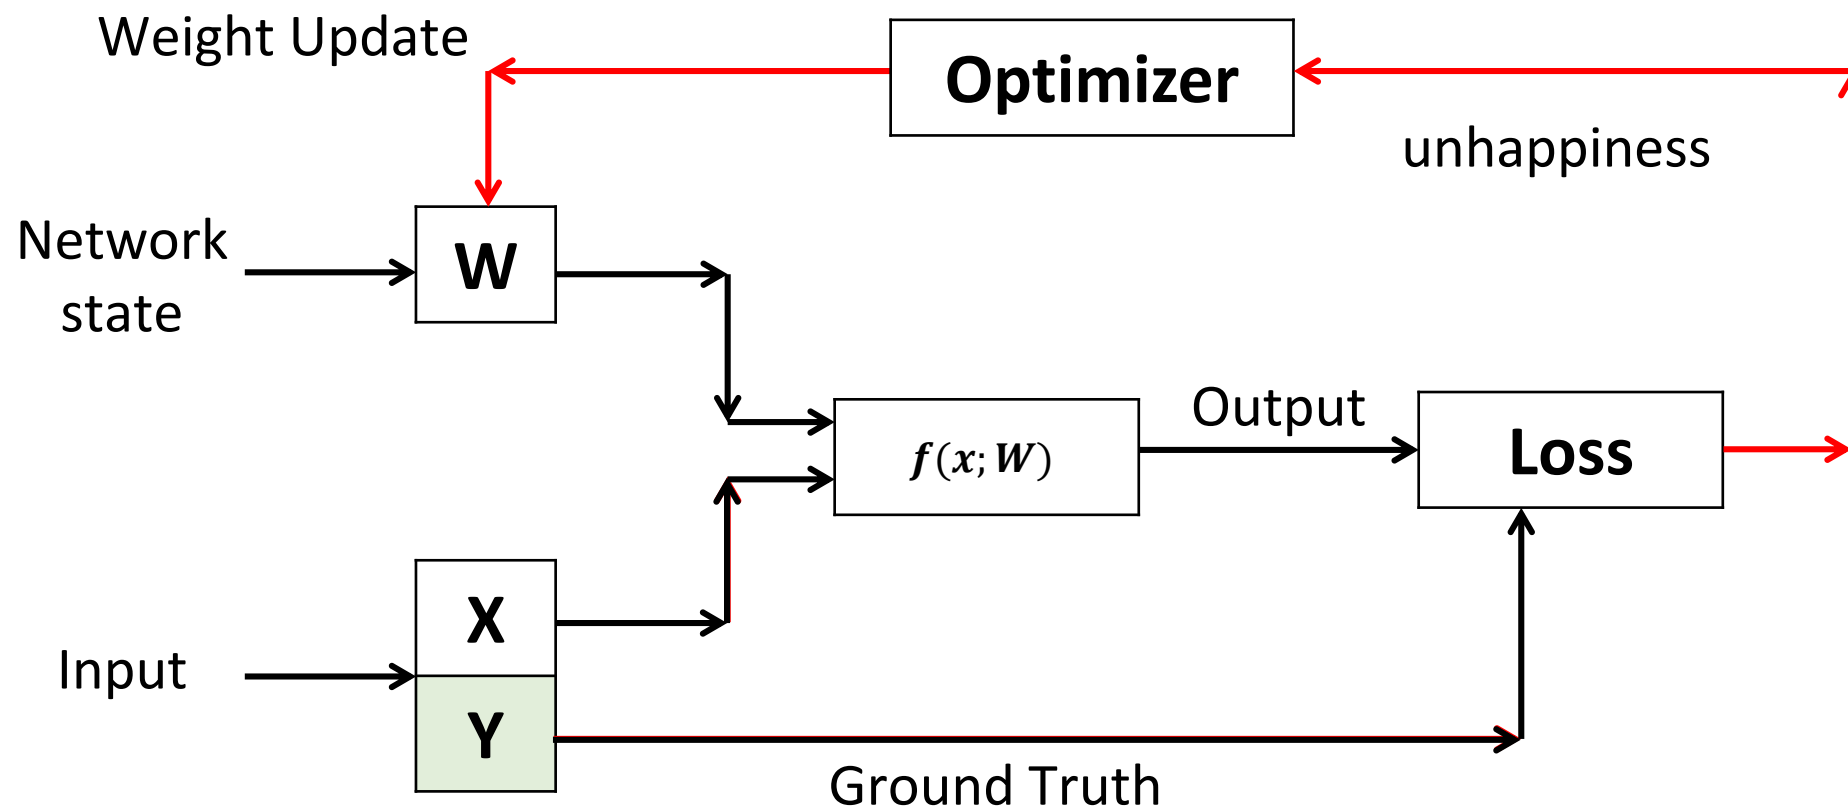
Evaluating

Utilities









Weight Update:

$$w_{ji} = w_{ji} + \Delta w_{ji}$$

$$\Delta w_{ji} = -\eta \frac{\partial E}{\partial w_{ji}}$$

There are two main types of models available in Keras:

- The Sequential model

```
model = Sequential()  
model.add(Dense(32, input_dim=784))  
model.add(Activation('relu'))
```

There are two main types of models available in Keras:

- The Sequential model
- The Model class

```
from keras.layers import Input, Dense
from keras.models import Model

# This returns a tensor
inputs = Input(shape=(784,))

# a layer instance is callable on a tensor, and returns a tensor
x = Dense(64, activation='relu')(inputs)
x = Dense(64, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x)

# This creates a model that includes
# the Input layer and three Dense layers
model = Model(inputs=inputs, outputs=predictions)
model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
model.fit(data, labels) # starts training
```




Data
Preprocessing

Defining
Model

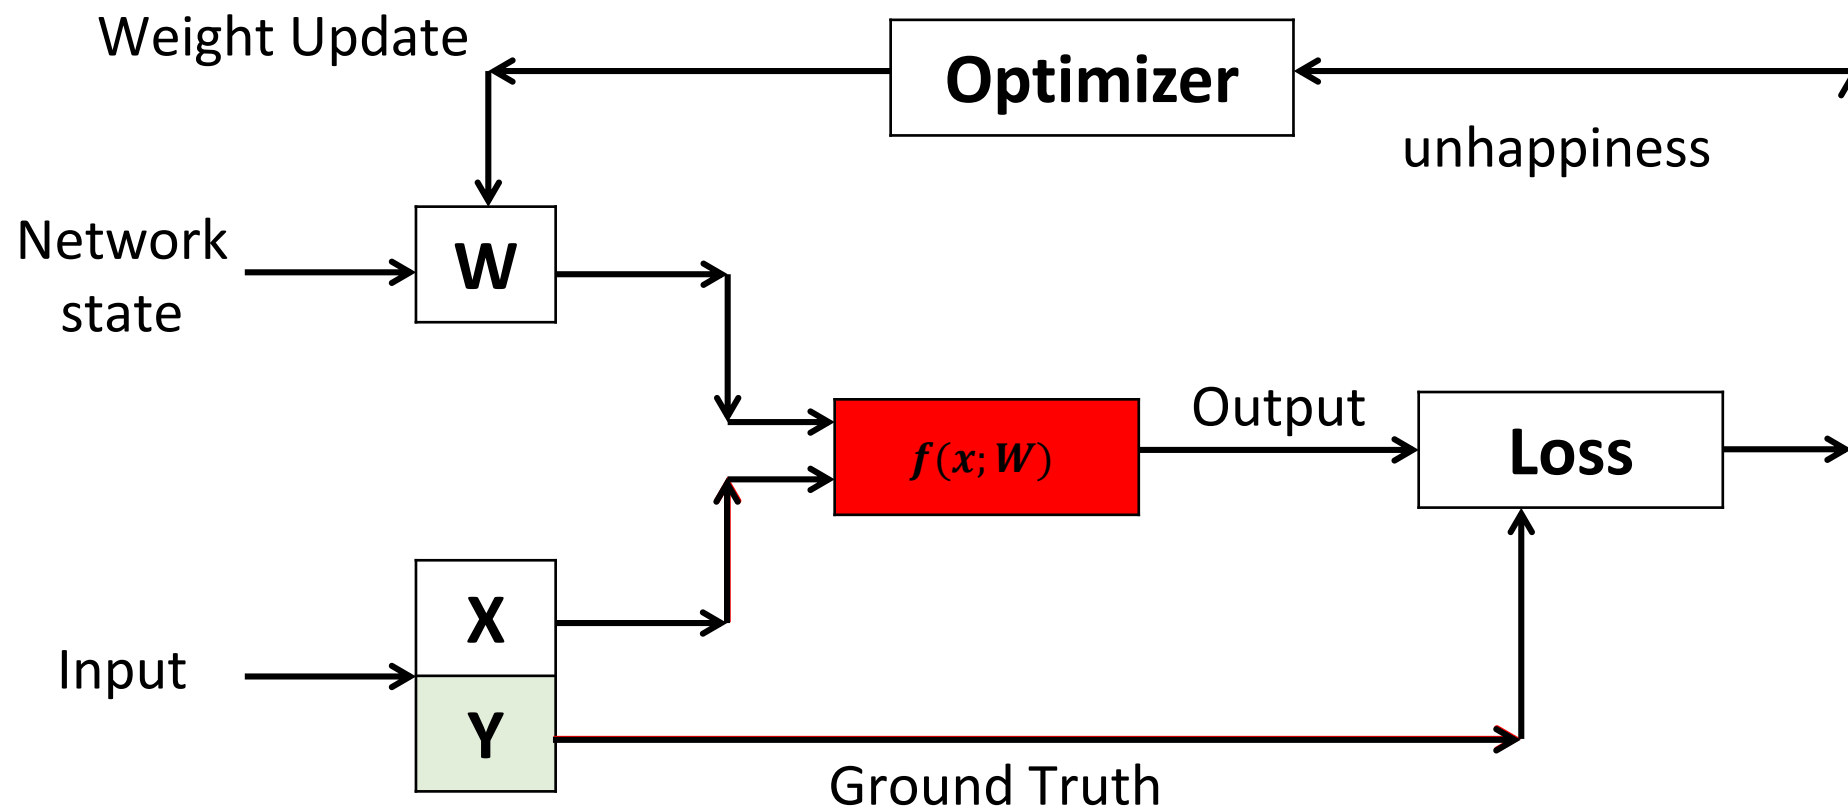
Training

Evaluating

Utilities


There are two main types of models available in Keras:

- The Sequential model
- The Model class
- Model subclassing



Keras.layers


- Core layers:
 - **Dense**
 - Activation
 - Input
 - Reshape
 - Flatten
 - ...
- Convolutional Layers (will be covered later)
- Recurrent Layers (will be covered later)



Keras.layers

- Core layers:
 - Dense
 - Activation
 - Input
 - Reshape
 - Flatten
 - ...
- Convolutional Layers (will be covered later)
- Recurrent Layers (will be covered later)


Sigmoid
Relu
Tanh
Softmax
...



Keras.layers

- Core layers:
 - Dense
 - **Activation**
 - Input
 - Reshape
 - Flatten
 - ...
- Convolutional Layers (will be covered later)
- Recurrent Layers (will be covered later)


Sigmoid
Relu
Tanh
Softmax
...



Keras.layers

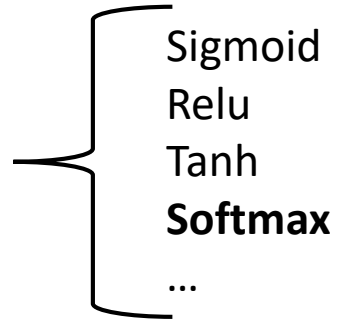
- Core layers:
 - Dense
 - **Activation**
 - Input
 - Reshape
 - Flatten
 - ...
- Convolutional Layers (will be covered later)
- Recurrent Layers (will be covered later)

Sigmoid
Relu
Tanh
Softmax
...



Keras.layers

- Core layers:
 - Dense
 - **Activation**
 - Input
 - Reshape
 - Flatten
 - ...
- Convolutional Layers (will be covered later)
- Recurrent Layers (will be covered later)



- Sigmoid
- Relu
- Tanh
- Softmax**
- ...

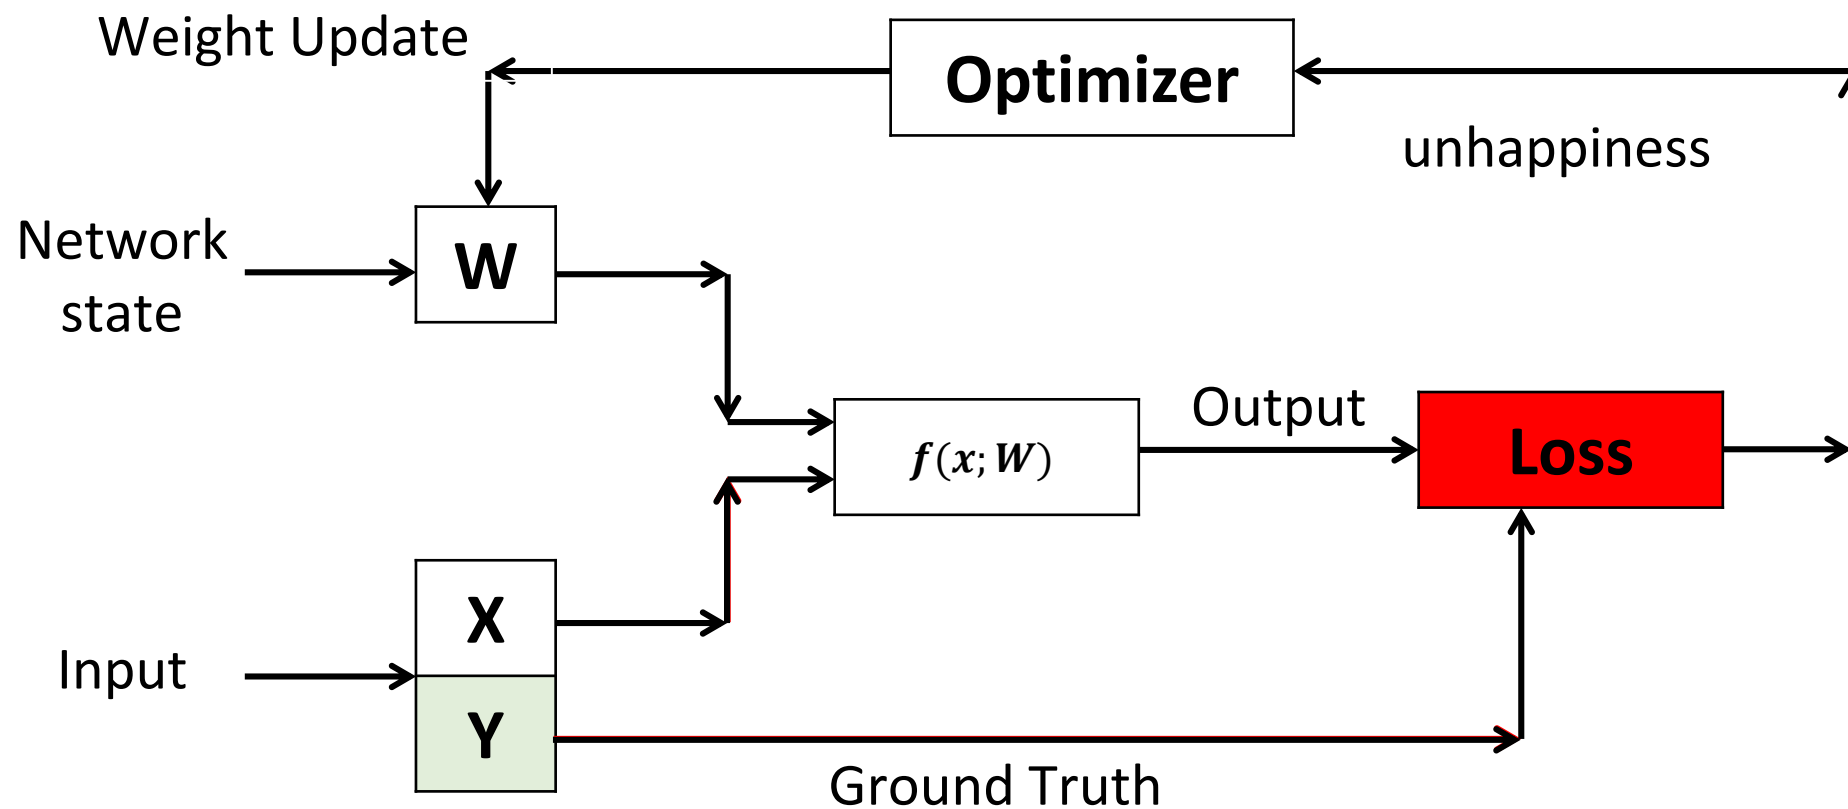
$$\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$$

Keras.layers

- Core layers:
 - Dense
 - Activation
 - **Input**
 - **Reshape**
 - **Flatten**
 - ...
- Convolutional Layers (will be covered later)
- Recurrent Layers (will be covered later)

Keras.layers

- Core layers:
 - Dense
 - Activation
 - Input
 - Reshape
 - Flatten
 - ...
- **Convolutional Layers (will be covered later)**
- **Recurrent Layers (will be covered later)**



Keras.losses

- **Mean_squared_error**
- Mean_absolute_error
- Categorical_crossentropy
- ...



$$MSE = \frac{1}{N} \sum_{i=1}^N (f_i - y_i)^2$$

$$MAE = \frac{1}{N} \sum_{i=1}^N |f_i - y_i|$$

$$H = - \sum_x y_i \log(f_i)$$

Keras.losses

- Mean_squared_error
- **Mean_absolute_error**
- Categorical_crossentropy
- ...



$$MSE = \frac{1}{N} \sum_{i=1}^N (f_i - y_i)^2$$

$$MAE = \frac{1}{N} \sum_{i=1}^N |f_i - y_i|$$

$$H = - \sum_x y_i \log(f_i)$$

Keras.losses

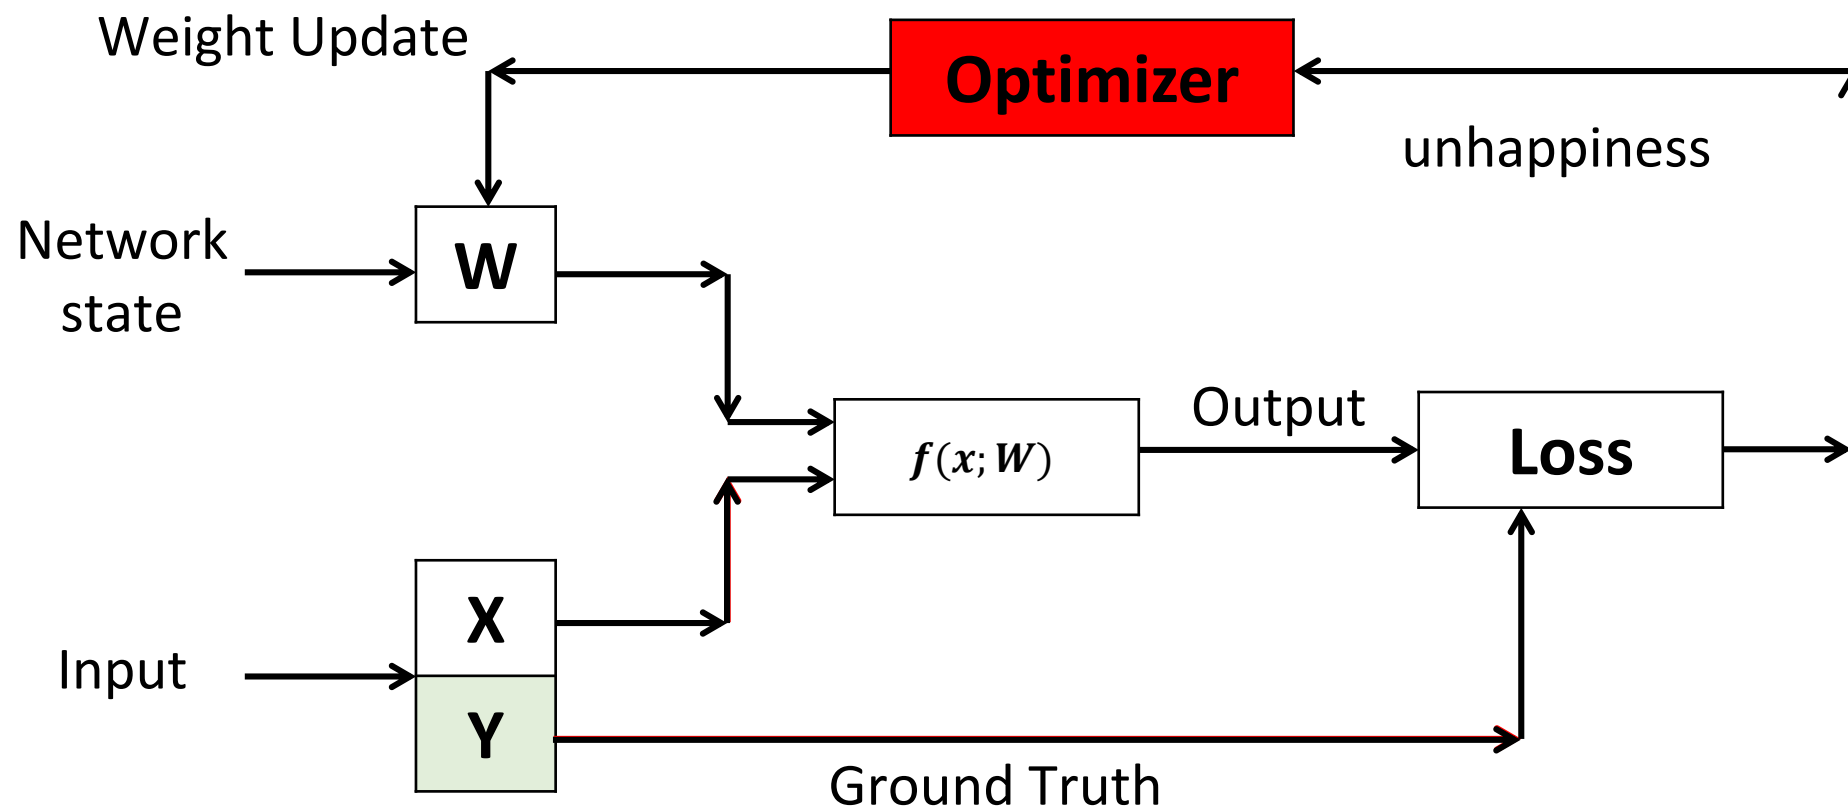
- Mean_squared_error
- Mean_absolute_error
- **Categorical_crossentropy**
- ...

$$MSE = \frac{1}{N} \sum_{i=1}^N (f_i - y_i)^2$$

$$MAE = \frac{1}{N} \sum_{i=1}^N |f_i - y_i|$$



$$H = - \sum_x y_i \log(f_i)$$




Data
Preprocessing

Defining
Model

Training


Evaluating

Utilities




Keras.Optimizers

- Stochastic Gradient Descent
- RMSprop
- Adam




Keras.initializer

- **RandomNormal**
- RandomUniform
- Glorot_uniform
- Glorot_normal



Keras.initializer


- RandomNormal
- **RandomUniform**
- Glorot_uniform
- Glorot_normal



Keras.initializer

- RandomNormal
- RandomUniform
- **Glorot_normal**
- Glorot_uniform


$$\sigma = \sqrt{\frac{2}{fan_{in} + fan_{out}}}$$

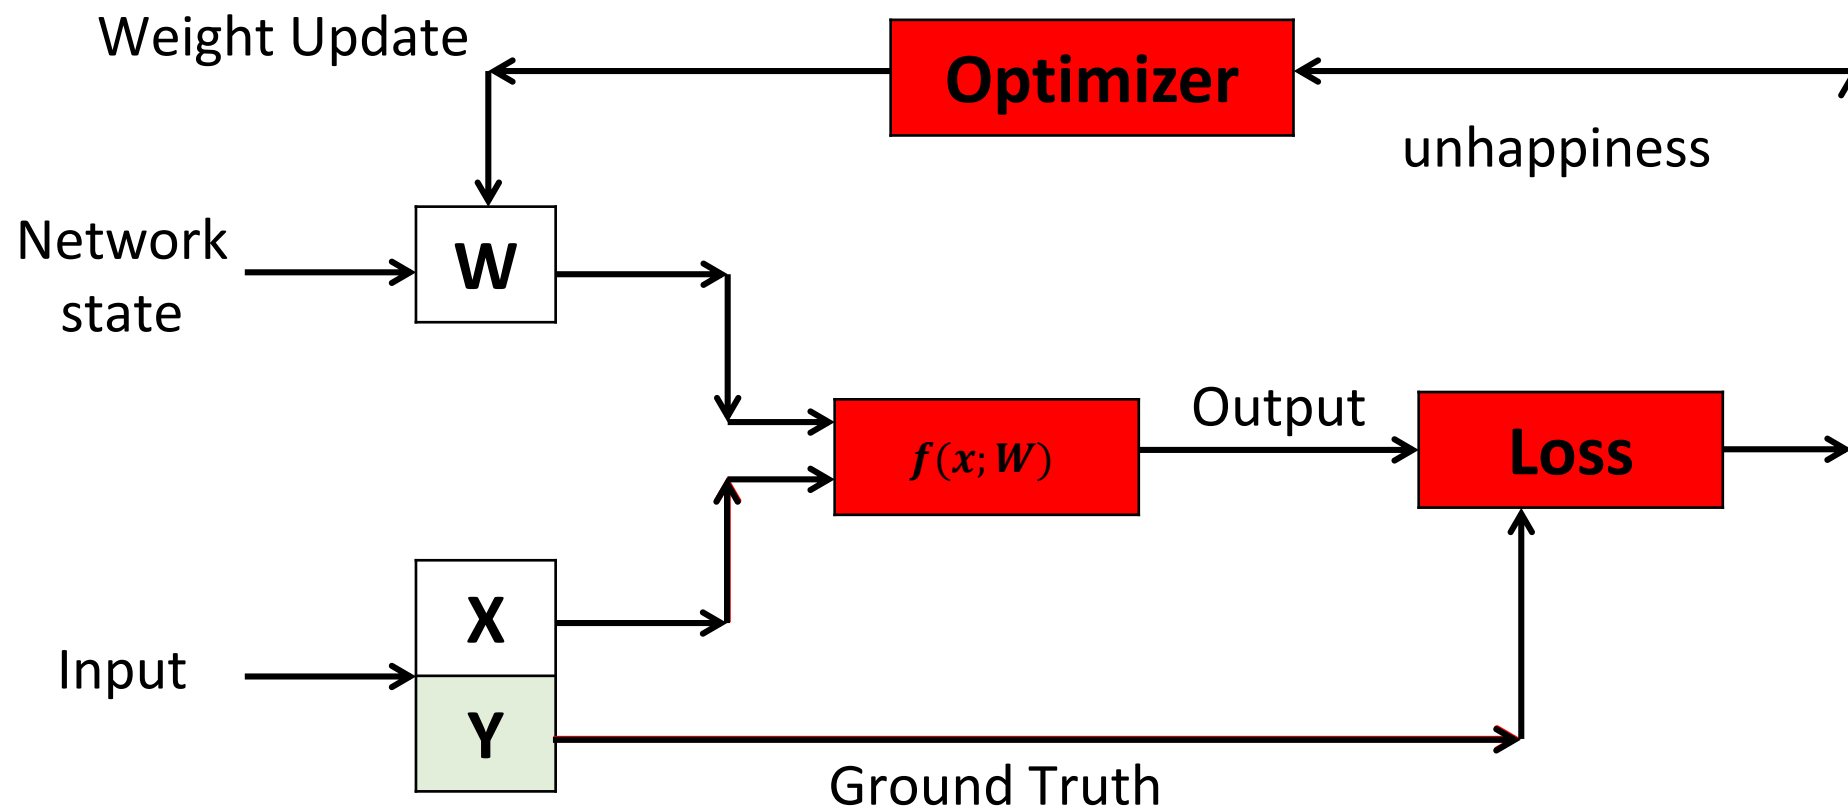


Keras.initializer

- RandomNormal
- RandomUniform
- Glorot_normal
- **Glorot_uniform**

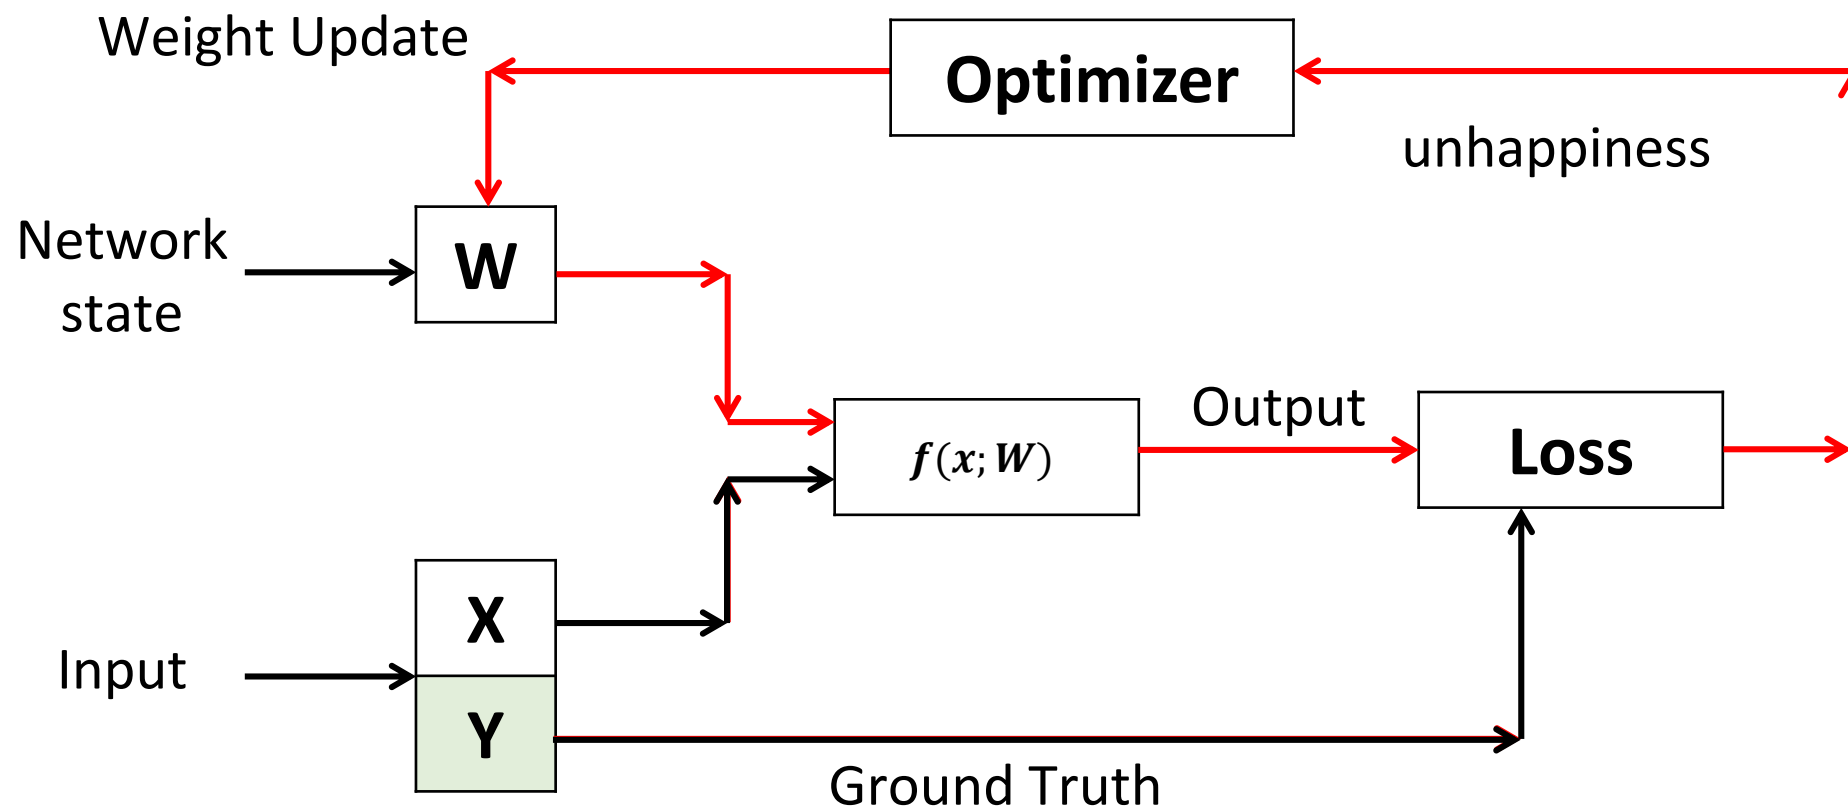
$$n = \sqrt{\frac{6}{fan_{in} + fan_{out}}}$$





model.compile

```
sgd = SGD(lr=0.01)
model.compile(loss='categorical_crossentropy',
              optimizer=sgd,
              metrics = ['accuracy'])
```


Data
Preprocessing

Defining
Model

Training

Evaluating

Utilities



Data
Preprocessing



Defining
Model


Training

Evaluating



Utilities

Batch1





Batch2



model.fit

```
history = model.fit(x_train, y_train,  
                    batch_size = 32,  
                    epochs = 3,  
                    verbose = 1,  
                    validation_split = 0.2)
```

model.evaluate

```
te_score = model.evaluate(x_test, y_test, verbose = 0)
```

model.predict

```
predicted_label = model.predict(img, 1)
```

Save/Load model

```
model.save('mlp.h5')
```

```
model = load_model('mlp.h5')
```


Keras.callbacks:

- **Terminate on NaN.**
- Model Checkpoint
- EarlyStopping
- LearningRateScheduler
- TensorBoard
- CSVLogger

The text "NaN" is written in a red, serif font. A horizontal red line is drawn underneath the text, emphasizing it. The background is a light gray square.


Keras.callbacks:

- Terminate on NaN.
- **Model Checkpoint**
- EarlyStopping
- LearningRateScheduler
- TensorBoard
- CSVLogger



Keras.callbacks:

- Terminate on NaN.
- Model Checkpoint
- **EarlyStopping**
- LearningRateScheduler
- TensorBoard
- CSVLogger



Keras.callbacks:

- Terminate on NaN.
- Model Checkpoint
- EarlyStopping
- **LearningRateScheduler**
- TensorBoard
- CSVLogger



Andrej Karpathy ✓
@karpathy



3e-4 is the best learning rate for Adam, hands down.

7:31 AM - Nov 24, 2016

♡ 397 💬 120 people are talking about this



Andrej Karpathy ✓
@karpathy

(i just wanted to make sure that people understand that this is a joke...)


12:21 PM - Nov 24, 2016

♡ 95 👤 See Andrej Karpathy's other Tweets




Keras.callbacks:

- Terminate on NaN.
- Model Checkpoint
- EarlyStopping
- **LearningRateScheduler**
- TensorBoard
- CSVLogger



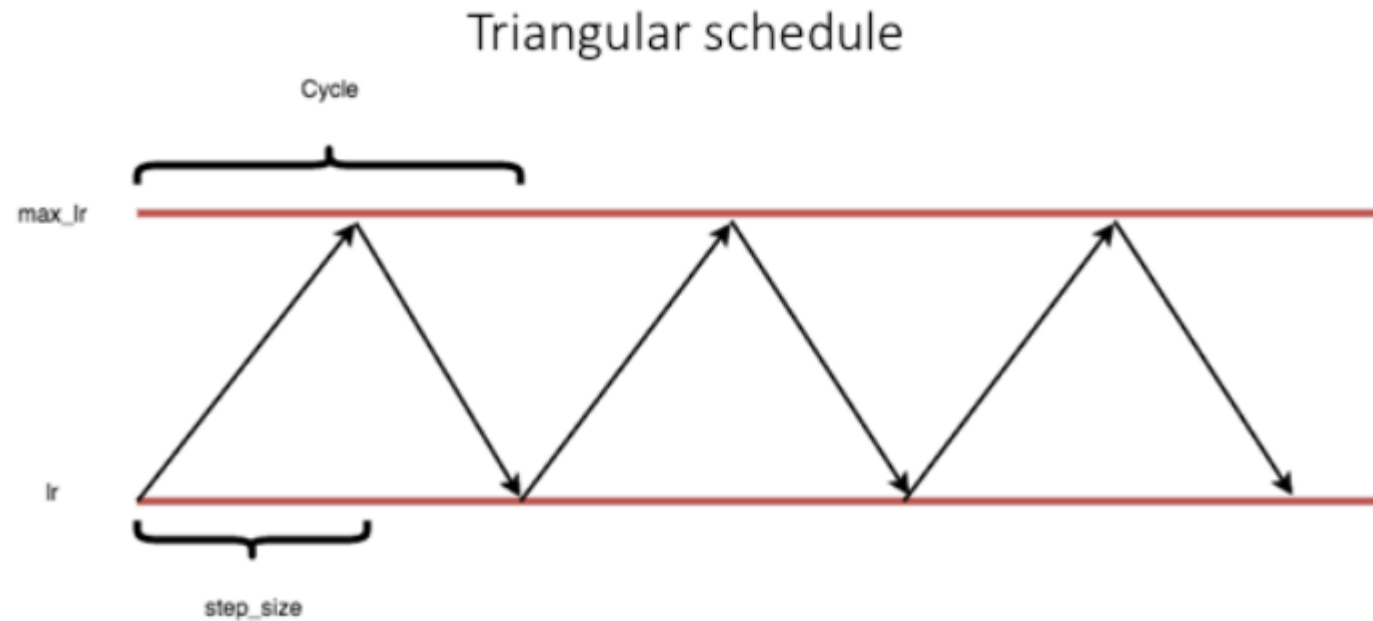
Keras.callbacks:

- Terminate on NaN.
- Model Checkpoint
- EarlyStopping
- **LearningRateScheduler**
- TensorBoard
- CSVLogger



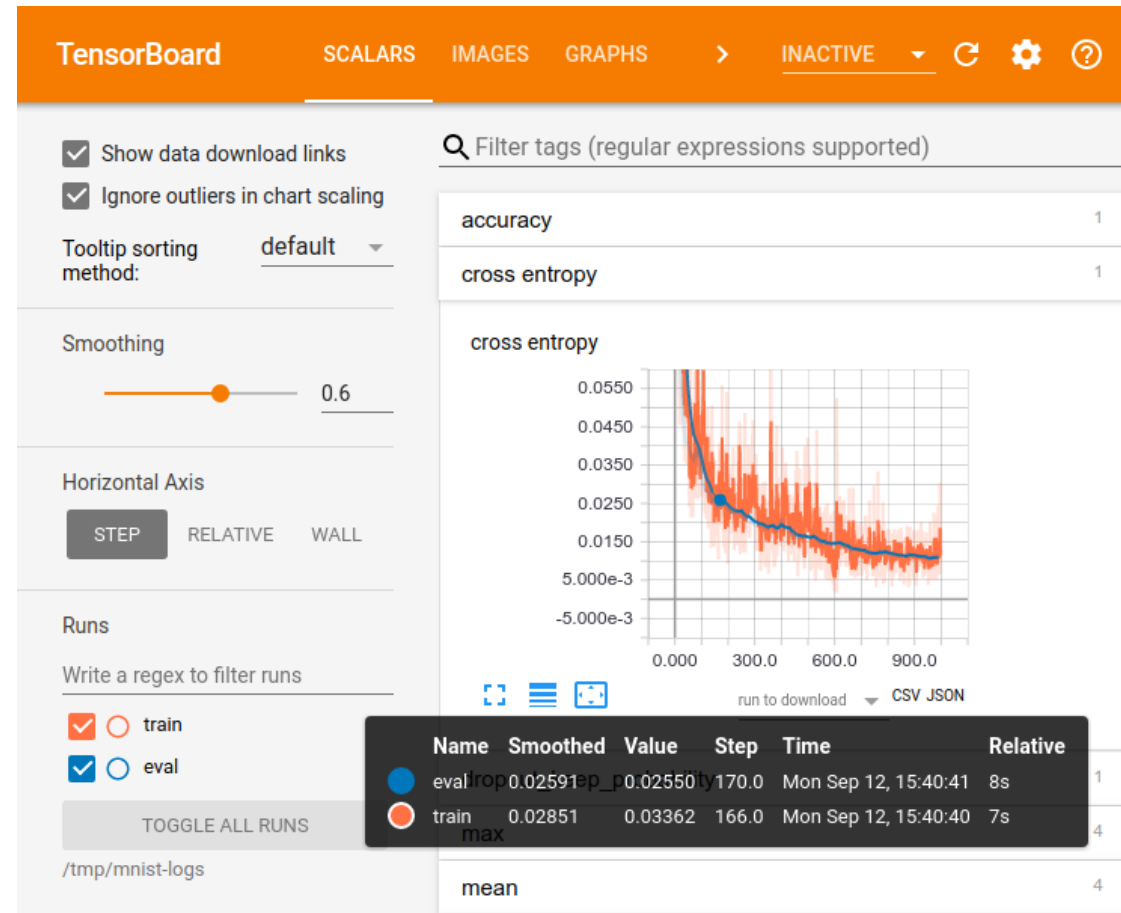
Keras.callbacks:

- Terminate on NaN.
- Model Checkpoint
- EarlyStopping
- **LearningRateScheduler**
- TensorBoard
- CSVLogger



Keras.callbacks:

- Terminate on NaN.
- Model Checkpoint
- EarlyStopping
- LearningRateScheduler
- **TensorBoard**
- CSVLogger




Keras.callbacks:

- Terminate on NaN.
- Model Checkpoint
- EarlyStopping
- LearningRateScheduler
- TensorBoard
- **CSVLogger**



Hands-On Session:

Training a fully-connected neural network on MNIST.



Thank You For Your Attention